

- Identify opportunities to improve processes and strategies with technology solutions

We Believe That You Have

- Minimum 3+ years of experience in Data Science
- Strong Python programming skills
- Solid SQL knowledge and experience with relational databases
- Experience in data wrangling and working with multi-source datasets
- Strong foundation in statistics, mathematics, and machine learning
- Fluency in both Hungarian and English
- Bachelor's or Master's degree in Computer Science, Statistics, Informatics, Information Systems, Physics, Mathematics, or another quantitative field

Other Preferred Experiences

- CRM and consumer journey analytics experience
- Cloud platform knowledge (AWS, Snowflake)
- Retail/store-level analytics experience

We offer you an attractive base salary of at least HUF 15.116.049 / year gross, with the willingness to pay more in line with the market depending on your



Senior Data Scientist

EPAM Systems • Budapest, Budapest, Hungary (...)



Requirements

- 3+ years of experience in data mining, statistics or machine learning
- Strong problem-solving skills and a data-driven mindset
- Solid foundation in applied mathematics, statistical analysis, and machine learning
- Excellent communication and presentation skills, including working proficiency in English
- Proficiency in SQL and relational databases
- Programming experience, preferably in Python
- Hands-on experience with data analysis tools and libraries (e.g., NumPy, SciPy, scikit-learn, pandas, matplotlib, R, SAS, SPSS, MATLAB)
- Experience with Big Data technologies (e.g., Spark, MLlib)
- Proficiency with at least one major cloud platform
- Experience deploying data science solutions to production environments
- Strong data visualization skills
- Bachelor's or Master's degree in Computer Science, Mathematics, Applied Statistics or a related field
- Strong written and verbal English communication skills (B2+)

Jelentkezés

Mentés

**The knowledge I own:****Your Profile**

What you should have.

- Academic background in data science, computer science, engineering, or statistics
- Proven skills in data science and strong knowledge of Python (especially PySpark), as well as experience writing complex SQL queries
- Mastery of machine learning to solve clustering, classification, regression, anomaly detection, simulation, and optimization problems using large datasets
- Demonstrated ability to combine diverse internal and external datasets and use them in a wide range of analyses and modelling tasks
- Experience with big-data technologies (preferably Azure and/or Databricks) and with data visualization tools, preferably e.g., Tableau

What is considered desirable

- Experience with retail data models and processes (especially in the areas of purchasing, marketing and/or planning)
- Experience with long-term analysis and demand/sales forecasting
- Experience with data science use cases in the supply chain or in finance
- Experience working with customer data from mixed systems

The offer that would convince me:

- Permanent contract with attractive compensation package
- High degree of personal responsibility and varied activities from day one
- International environment with colleagues from all over the world
- IT projects that set new standards in the retail industry

Riasztás beállítása hasonló állásokhoz

Senior Data Scientist, Budapest, Hungary

Ki

**Toborzással kapcsolatos betekintés
zárolásának feloldása itt: ALDI
International IT Services (Hungary)**

Jelentkezés

Mentés

Your skills will include (must have)

- 5+ years of professional experience in data science, quantitative research, or statistical modelling
- Deep, intuitive understanding of probability, statistics, and machine learning theory
- Expert-level skills in the Python data stack. You write clean, production-grade code
- Proven experience designing complex backtesting environments and defining custom evaluation metrics for unique business problems

Nice to have

- Strong understanding of betting concepts, odds, overround, market-making, or quantitative trading. You understand what it means to build models where calibrated probabilities directly drive revenue
- Experience modelling off streaming data or data that updates continuously over time
- Deep knowledge of competitive esports (CS2, Dota 2, LoL), the underlying game mechanics, and the competitive meta
- Understanding of modern MLOps principles and experience with tools like MLFlow, Airflow, etc.

Téged keresünk, ha rendelkezel az alábbiakkal:

- **Legalább 3 év** releváns tapasztalat adatmérnök vagy adattudós szerepkörben
- Magabiztos **Python és SQL** tudás
- Tapasztalat **adatmodellezésben és adatintegrációban**
- Jártasság **statisztikai elemzésben, adatmodellezésben és gépi tanulásban**
- Fejlett **analitikus gondolkodásmód és problémamegoldó készség**

Előnyt jelent, ha vannak:

- SAS ismeretek
- Befejezett felsőfokú tanulmányok matematikai, statisztikai, mérnöki, informatikai vagy gazdasági területen
- Kiváló kommunikációs készség

General qualifications:

- Fluent English speaker with excellent communication skills, comfortable in client-facing roles.
- 5 years of experience in data science or a related field.
- Hands-on experience with LLM projects, machine learning, deep learning.
- Affinity to building product / product-oriented vision
- Domain expertise in any of the following fields is highly preferred: international finance, healthcare, pharma or meteorology.

Technical requirements:

- Strong Python and SQL skills
- Leverage cloud-based data science tools on AWS (SageMaker, Bedrock), Azure, or GCP for scalable model training and deployment
- IDE (e.g.: Jupyter, VSCode)
- Version control Git / BitBucket / AZ DevOps
- Core machine learning libraries (sklearn, LightGBM, torch)
- Traditional ML projects (supervised and unsupervised)
- NLP experience / affinity: Understand concepts of text classification, information extraction
- Generative AI & LLM frameworks - HuggingFace / transformers, OpenAI; proprietary API & open-source models
- Understand concepts of RAG, build & maintain pipeline, understand concepts of pretraining, fine-tuning

Nice to have:

- Databricks / distributed computing / scalable ML
- Familiarity with database systems (SQL, NoSQL)
- ML Lifecycle management (e.g.: MLFlow / wandb)
- Web frameworks (e.g.: Streamlit / FastAPI / Flask)
- Automatization pipelines on cloud (e.g.: Azure Functions, AWS Lambda, Databricks notebook jobs)
- API integration
- Experience with deploying ML models in production environments
- Experience with LangChain / LlamaIndex
- Understanding of LLM Agents, agentic behavior, prompt engineering
- Being able to differentiate what 'needs' LLMs and what can be solved with traditional ML / NLP

Additional Information

Requirements

We are looking for smart people with good communication skills who are hungry to tackle interesting challenges. If the responsibilities are attractive to you and you are passionate about developing novel models and improving on the state of the art, we would love to hear about you.

We are impressed if you...

- Have a Master's or Ph.D. in Computer Science, Computer Engineering, Computer Vision, or a related field
- Have experience in state-of-the-art Generative AI technologies
- Have experience designing and implementing AI solutions using Python
- Are proficient in ML and CV toolkits, e.g., PyTorch, openCV
- Experience using cloud computing services
- Have experience transferring technology from research into a shipping product
- Are able to present your results to both technical and non-technical audiences

The soft skills that we think are important

- Clearly explain problems, present solutions, and debate strategies in English to both technical and non-technical audiences
- Understand a business problem, then propose and weigh multiple solutions
- Collaborate effectively with Data Scientists, Engineers, and Product Team members
- Estimate and plan machine learning projects and coordinate with engineering team members

Bonus points

- Experience with A/B testing as a form of model evaluation
- Hands-on experience with PySpark and/or Databricks
- Experience handling unstructured text and/or NLP methodology
- Experience building tooling, scripts (e.g. bash, python, etc), or systems to support ML investigations, deployments, etc.
- Monitoring production model performance using automated dashboards or data visualization
- You are active on GitHub and have contributed to open-source projects

If the role sounds interesting, apply now and get to know us during the interviews. You can read more about our hiring process on Glassdoor.

First name*

Last name*

Jelentkezés [↗](#)

Mentés

Required

- professional experience in data science / ML, with building LLM or GenAI applications in production.
- Hands-on experience with agentic frameworks LangChain or LangGraph.
- Strong understanding of LLM application patterns including prompt engineering
- Expert-level Python (pandas, NumPy, scikit-learn, async programming, API development with FastAPI).
- Experience with vector databases and embeddings (Azure AI Search, Pinecone, Weaviate, Qdrant, FAISS, or pgvector).
- Strong SQL skills and experience working with large-scale data.

Preferred

- Bachelor's or Master's degree in Computer Science, Data Science, Engineering, Mathematics, or a related field.
- Hands-on experience with Azure services: Azure OpenAI / AI Foundry, Azure ML, Azure AI Search, Azure Databricks, Azure Data Factory, or Synapse.
- Solid grounding in classical ML: supervised/unsupervised learning, model evaluation, and hyperparameter tuning; experience with TensorFlow or PyTorch.
- Experience with Model Context Protocol (MCP), OpenAI Assistants/Agents SDK, or Anthropic tool-use APIs.
- Strong understanding of function/tool calling, structured outputs, RAG, agent memory, and multi-agent orchestration.
- Proven MLOps/LLMOps experience: MLflow, prompt/model versioning, CI/CD (Azure DevOps or GitHub Actions), and production monitoring.
- Ability to evaluate and mitigate LLM-specific risks: hallucination, prompt injection, data leakage, and cost/latency constraints.
- Experience with AutoGen, CrewAI, Semantic Kernel.
- Microsoft certifications: AI-102 (Azure AI Engineer), DP-100 (Azure Data Scientist Associate).
- Experience fine-tuning open-source LLMs (Llama, Mistral, Phi) using LoRA/QLoRA and serving via vLLM or Azure ML endpoints.
- Familiarity with observability/tracing for agents: LangSmith, Langfuse, Arize Phoenix, or OpenTelemetry.
- Knowledge of containerization and deployment: Docker, Kubernetes (AKS), Azure Container Apps.
- Big data experience with Apache Spark (PySpark) via Azure Databricks.
- Experience with knowledge graphs, graph RAG, or semantic layers for agent grounding.
- Contributions to open-source GenAI/agentic projects or published technical content.

We are looking for you if have:

- a completed Master's degree (ideally in Statistics, Mathematics or Computer Science) or up to 3 years of relevant work experience,
- proven experience using Python and basic knowledge of shell scripting,
- familiarity with software development methods and version control tools such as Git,
- ability to work independently and reliably in a small team,
- proactivity and eagerness to learn,
- willingness to experiment with various generative and descriptive AI use cases,
- at least intermediate proficiency in English (oral and written); Hungarian knowledge is an advantage.

We consider it an advantage if you have:

- experience with database technologies (MS SQL, PostgreSQL, MongoDB),
- experience with deep learning models and related Python packages such as Transformers or PyTorch,
- experience in developing NLP models,
- experience with Azure.



Lead Data Scientist



Sony Europe • Budapest, Budapest, Hungary (H...

practices to continually improve AI capabilities
and delivery excellence.

You Will Bring

- Proven experience leading Data Science, Machine Learning, or AI teams within a large and complex organisation.
- Strong expertise in Machine Learning, Generative AI, and modern AI technologies, e.g. AI Agents, with the ability to identify high-impact business applications.
- Experience delivering AI solutions from concept through production in enterprise environments.
- Strong leadership skills with experience leading and developing technical teams.
- Hands-on experience with Python, SQL, and modern AI/ML frameworks and libraries.
- Understanding of MLOps practices, data pipelines, model deployment, and operationalisation of AI solutions.
- Strong stakeholder management and communication skills, with the ability to bridge business and technical audiences.
- Fluent English.

Your new office

Jelentkezés 

Mentés



You Have

- Degree in a quantitative field such as Computer Science, Mathematics, Statistics, Physics, or related discipline.
- Solid problem solving and mathematical foundations, and an interest in the financial industry.
- Strong communication skills (ability to present complex ideas clearly).
- Strong Python skills and experience building applications and backend services.
- Experience with data analysis and databases (SQL / time-series).
- Familiarity with LLM / GenAI patterns (RAG, tool use, AI-assisted coding).
- Analytical mindset, high attention to detail, and ability to handle multiple projects.

Nice To Have

- Experience tuning prompts and interactions to improve LLM behaviour and UX.
- Hands-on GenAI work with LangChain, LangGraph or similar frameworks.
- Understanding of neural networks and modern NLP / LLM concepts.
- Experience designing tests and evaluations for GenAI systems.
- Experience taking projects from prototype to production.
- Experience with RAG, structured generation (JSON / SQL / function calling), and agent / tool-calling setups.
- Experience working in an agile team.

#BPMM

What You Can Expect From Morgan Stanley

At Morgan Stanley, we raise, manage and allocate

[Jelentkezés](#)[Mentés](#)

Requirements

- **Experience:** 5+ years of hands-on experience in data science and machine learning, with a proven portfolio of impactful projects.
- **Modelling Expertise:** Experience in training tabular and time-series models using classical machine learning and deep learning approaches.
- **Technical Skills:** Proficiency in building data pipelines, training models, and conducting experiments using industry-standard tools.
- **Database Knowledge:** Advanced SQL skills, including the ability to debug complex queries and a solid understanding of OLTP vs. OLAP systems (experience with OLAP systems is a plus).
- **Problem-Solving:** Exceptional analytical skills with a detail-oriented and innovative approach to challenges.
- **Communication:** Strong verbal and written communication skills to collaborate effectively within a team.

Bonus points

- **Cloud & Containers:** Experience with cloud platforms (AWS, GCP, Azure) and container tools like Kubernetes.
- **ML Pipelines:** Expertise in ML pipeline development, including data preprocessing, model training, deployment, and monitoring.
- **MLOps Tools:** Familiarity with MLOps tools like MLflow or Feast.
- **OLAP Expertise:** Hands-on experience with large-scale OLAP systems such as Clickhouse, BigQuery, or Snowflake.
- **Advanced Modeling:** Knowledge of advanced techniques for handling imbalanced datasets and other complex modeling challenges.

Ezek a feladatok várnak nálunk:

Analitikai és adattudományi feladatok

- Prediktív modellek és előrejelzések készítése
- Ügyféljövödelmezősségi és értékteremtési modellek
- Statisztikai elemzések és üzleti hipotézisek validálása
- Szenárióelemzések készítése a stratégiai döntések támogatására
- Machine Learning megoldások bevezetése.

Adatkezelési és riorting feladatok

- Nagy mennyiségű adat feldolgozása és elemzése
- Adatminősségi ellenőrzések és validációk, ezekhez automatizált megoldások készítése
- Power BI riortok fejlesztése és karbantartása

Döntéstámogatási feladatok

- Vezetői elemzőanyagok készítése (Prezentációk és döntéselőkészítő anyagok)
- Üzleti területekkel és leánybankokkal való együttműködés
- KPI-rendszerek kialakítása és fejlesztése (Új mutatók, benchmarkok és teljesítménymérési megoldások kidolgozása)

Téged várunk, ha van

- Felsőfokú végzettséged gazdasági, matematikai, statisztikai, pénzügyi, mérnöki vagy informatikai területen
- Legalább 2-3 év hasonló munkakörben szerzett gyakorlatod
- SQL magas szintű ismereted
- Python vagy R programozási ismereted
- Adatmodellezési tapasztalatod
- Magabiztos angol nyelvtudásod
- Kiváló kommunikációs és prezentációs készséged

Előnyvel indulsz nálunk, ha van

- Ügyféljövödelmezősségi vagy értékteremtési modellek fejlesztésében tapasztalatod
- Banki, pénzügyi vagy controlling területen szerzett tapasztalatod
- Power BI fejlesztési tapasztalatod
- Machine Learning projekteken szerzett ismereteid
- Databricks vagy felhős adatplatformok ismereted
- Predictive Analytics vagy AI megoldások gyakorlati alkalmazásában gyakorlatod

We are interested in professionals with expertise in

- Artificial Intelligence
- Machine Learning
- Computer Vision
- Deep Learning
- Data Science
- Applied Research
- Related technical fields

What makes you a great fit

- A strong foundation in AI, Machine Learning, Computer Vision, or related disciplines
- Passion for solving complex problems through technology and innovation
- Strong analytical thinking and problem-solving abilities
- Curiosity, adaptability, and a continuous learning mindset
- Ability to work independently while collaborating effectively with cross-functional teams
- A hands-on approach and willingness to take ownership of challenges
- Excellent communication and collaboration skills
- Enthusiasm for working in a fast-paced startup environment where experimentation and innovation are encouraged

Bonus

- Experience working in a fast-paced startup or scale-up environment
- Experience deploying AI solutions into production environments
- Experience with retail, food-tech, or automation technologies

How you'll have an impact:

- Research emerging fraud patterns in real-time with our Risk Intelligence team
- Improve the important components of the Signifyd Commerce Protection Platform
- Communicate complex ideas to a variety of audiences, including executives
- Build production machine learning models that identify fraud
- Write production and offline code in python, PySpark
- Work with distributed data pipelines
- Collaborate with engineering teams to strengthen our machine-learning pipeline

Past experience you'll need:

- A degree in computer science or a comparable analytical field
- 3+ years of post-undergrad work experience required
- Strong verbal and written communication skills
- Strong machine learning and statistical background, and a track record of being able to deliver under pressure.
- Write code and review others' in a shared codebase in Python
- Practical SQL knowledge
- Design experiments and collect data
- Familiarity with the Linux command line
- Candidates must be based in Hungary

Bonus points if you have:

- Previous work in fraud, payments, or e-commerce
- Data analysis in a distributed environment
- Passion for writing well-tested production-grade code
- A Master's Degree or PhD



Requirements

- 4+ years of experience in AI/ML research or applied machine learning
- Strong background in mathematics and computer science, including statistics, optimization, linear algebra, and algorithms
- Solid understanding of machine learning and deep learning methods
- Experience working with real-world datasets, including data exploration, cleaning, feature extraction, and experimental validation
- Proficiency in Python and common scientific/ML libraries (NumPy, SciPy, Pandas, etc.)
- Experience with deep learning frameworks such as TensorFlow or PyTorch
- Familiarity with computer vision or multimodal AI (images, video, 3D data such as point clouds and meshes)
- Ability to independently formulate research hypotheses, design experiments, and evaluate results
- Strong analytical thinking and ability to solve complex problems with limited prior solutions

Nice to have

- Experience with 3D data processing (point clouds, meshes, spatial data)
- Experience building or experimenting with AI agents or LLM-based systems
- Experience training and fine-tuning neural networks
- Interest in geometry, graph theory, or computational geometry

What we offer:

[Jelentkezés](#)[Mentés](#)

YOUR PROFILE

- Education level / Academic background
- MSc student or recent graduate in Artificial Intelligence, Computer Science, Data Science, Informatics, Engineering, or a related quantitative/technical field.

Hard Skills

- Strong English proficiency
- Python for AI & prototyping (pandas, scikitlearn, basic model evaluation)
- SQL & data wrangling (exploring, cleaning, shaping structured/unstructured data)
- GenAI prototyping & LLM fundamentals (prompting, RAG basics, understanding hallucination risks)
- Rapid prototyping mindset (combining Python, APIs, Copilot Studio, simple UIs)
- Power BI & visualization basics (exploratory dashboards, clear insight presentation)
- Securebydesign & responsible AI awareness (data protection, privacy, ethical AI practices)
- Documentation & reproducibility discipline (clear assumptions, data flows, prompts, validation steps)

Soft Skills

- Collaboration & teamwork
- Problem-solving
- Initiative & proactivity
- Learning agility
- Resilience & ability to deal with feedback
- Time management & prioritization

If you're a recent graduate ready to commit to a meaningful 2-year journey, this opportunity is exactly what you're looking for. The 1-year international assignment gives the chance to build a long-term career with us - with the ambition to consider successful participants for future opportunities within Ceva.



Required Qualifications

- 6+ years of experience in data science, applied machine learning, or a closely related role.
- Strong mathematical, statistical, and machine learning foundations, including probability, statistics, optimization, and model evaluation.
- Proven ability to select, apply, and critically evaluate ML models and algorithms for real-world problems.
- Strong Python skills for analysis, modelling, experimentation, and prototyping.
- Strong SQL skills for data exploration, transformation, and analytical workflows.
- Excellent analytical thinking and problem-structuring abilities; comfort operating independently with loosely defined goals.
- Experience using Git for version control and collaborative development.
- Strong English communication skills, both written and verbal.

Preferred Qualifications (Strong Plus)

- Hands-on experience with LLMs, including prompt/system design and building real-world applications.
- Experience with RAG systems, including retrieval strategies, chunking, evaluation, and performance tuning.
- Experience designing or contributing to agent-style AI systems and familiarity with concepts such as agent evaluation, guardrails, and reliability testing.
- ML modeling experience (e.g. supervised learning, ranking, classification) beyond exploratory analysis.
- Understanding of software engineering best practices, including testing strategies and CI/CD concepts.
- Experience working in Azure or similar cloud environments.
- Familiarity with Snowflake (and optionally Snowflake AI) as part of a modern data stack.
- Experience collaborating closely with domain experts; financial domain exposure is a plus but not required for strong technical candidates.

What Success Looks Like In This Role

- You independently deliver high-quality AI prototypes and evaluation frameworks for a defined subdomain or application module.
- You proactively define scope, success metrics,

[Jelentkezés](#)[Mentés](#)

Your Qualifications, Your Influence

To be successful in this role, you should have:

- PhD in Computer Science, Computer Vision, Machine Learning, Artificial Intelligence, Applied Mathematics, Electrical Engineering, Robotics, or closely related technical field.
- 7+ years of hands-on experience in Machine Learning / Deep Learning, with strong focus on Computer Vision.
- Strong practical experience with PyTorch and / or TensorFlow.
- Strong Python development skills.
- Experience with OpenCV, NumPy, Pandas, scikit-learn, and modern Python ML ecosystem.
- Deep understanding of classical Computer Vision algorithms and modern deep learning approaches.
- Strong experience with object detection, semantic segmentation, image classification, feature matching, image retrieval, and model evaluation.
- Experience with modern Computer Vision architectures and techniques, including CNNs, Transformers, Vision Transformers, YOLO, Mask R-CNN, CLIP-like models, SAM-like models, or similar.
- Experience bringing ML models into production environments.
- Experience with model optimization for inference speed, latency, memory usage, scalability, and reliability.
- Experience with REST APIs, Docker, CI / CD, model versioning, experiment tracking, and MLOps practices.
- Strong understanding of datasets, data quality, annotation processes, labeling requirements, and model error analysis.
- Ability to read, understand, and evaluate technical documentation and research papers in English.
- Ability to explain complex technical topics to engineering, product, and business stakeholders.
- Experience working in Agile software development environment.
- Strong ownership mindset, good judgment, and ability to make practical technical decisions under uncertainty.
- Comfortable working in distributed teams across multiple locations and time zones.

Preferred Skills and Technical Familiarity

The following experience will be helpful:

- Post-PhD research or industry experience in applied Computer Vision.
- Publications, patents, or strong applied research record in Computer Vision, Machine Learning, or AI.
- Experience leading technical direction for AI / ML projects.
- Experience mentoring ML engineers, software engineers, or data annotation teams.
- Experience with edge or mobile inference technologies, including ONNX, TensorRT, OpenVINO, TFLite, CoreML, or similar.
- Experience with large-scale image processing pipelines.
- Experience with synthetic data generation, active learning, weak supervision, or dataset quality improvement.
- Experience with multimodal models, vision-language models, prompt engineering, OpenAI, or similar AI platforms.
- Experience with cloud ML platforms and production monitoring of ML models.
- Familiarity with mobile applications, warehouse workflows, field operations systems, or image capture workflows.
- Familiarity with Azure DevOps, Git, CI / CD tooling, documentation systems, and practical software delivery processes.
- Experience in electronic components, technology distribution, supply chain, logistics, manufacturing, e-commerce, or similar B2B environments.

Success in the First 90 Days

Your Qualifications, Your Influence

To be successful in this role, you should have:

- PhD in Computer Science, Machine Learning, Artificial Intelligence, Computational Linguistics, Applied Mathematics, Data Science, or a closely related technical field
- 8+ years of professional experience in machine learning, artificial intelligence, or NLP
- 5+ years of hands-on experience building NLP models for production or near-production systems
- Deep understanding of modern neural network architectures, including RNN, CNN, Transformer-based architectures, attention mechanisms, embeddings, fine-tuning strategies, layers, modules, and loss functions
- Strong practical experience with NLP tasks such as NER, classification, text generation, semantic similarity, information extraction, and document understanding
- Strong experience with Large Language Models, including model evaluation, prompt design, fine-tuning, retrieval-augmented generation, and safe production usage
- Practical understanding of RAG, GraphRAG, knowledge graphs, embeddings, and hybrid retrieval approaches for production LLM applications
- Strong hands-on experience with Python
- Strong experience with PyTorch and Hugging Face Transformers
- Experience with ONNX or other model optimization / model serving formats
- Strong understanding of data preparation, data quality, labeling workflows, annotation guidelines, and model evaluation metrics
- Practical experience with main data analysis and machine learning libraries, including Pandas, NumPy, SciPy, scikit-learn, and Matplotlib
- Experience working with SQL databases and structured business data
- Experience with cloud platforms such as Microsoft Azure or AWS
- Ability to design experiments, define success metrics, compare model approaches, and explain trade-offs clearly
- Strong written and verbal English communication skills
- Experience working in Agile engineering environments
- Ability to provide technical leadership without requiring formal people management authority

Preferred Skills and Technical Familiarity

The following experience will be helpful:

- Experience leading NLP or AI research initiatives in a commercial production environment
- Experience with multilingual NLP systems
- Experience with vector databases, embeddings, semantic search, and RAG architectures
- Experience with knowledge graph concepts, including entity and relationship modeling, graph schema design, traversal queries, and LLM integration with graph databases such as Neo4j, FalkorDB, or similar technologies
- Experience with model serving, monitoring, drift detection, and production ML observability
- Experience with Docker and containerized ML workloads
- Experience with MLOps practices and CI/CD for machine learning systems
- Experience working with data annotation teams and creating annotation instructions
- Experience with .NET / C#, ASP.NET Core, or integration of ML services into enterprise software platforms
- Experience building prototypes, demos, and proof-of-concept applications for new AI capabilities
- Publications, patents, or recognized technical contributions in NLP, machine learning, or applied AI are a plus

Success in the First 90 Days

During the first 90 days, the Principal NLP Scientist is expected to:

- Understand the current NLP and AI automation landscape inside the company

Qualifications:

- MSc (PhD preferable) in computer science, statistics, computational biology or related fields advantageous
- At least 5 years of experience in data science / business intelligence / statistics role supporting AI/ML projects in clinical research, pharmaceutical, CRO or medical device setting
- Extensive programming experience in Python technology stack (pandas, numpy, scikit-learn, scikit-survival, matplotlib, seaborn, shap etc.) and experience in working with virtual environments and version control (e.g. git)
- Excellent knowledge of machine learning algorithms for clustering, classification, regression, anomaly detection, and optimization, preferably with both small and big data
- Experience in AI/ML applications in cardiology will be considered a plus
- Exposure to cloud computing platforms, preferably AWS and Google Cloud
- Excellent English language skills, storytelling and communications skills both oral and written, especially in explaining complex concepts in simple terms

Requirements:

Start: 05.10.2026

Duration: 12 months

Capacity: 5 days per week

Location: 100% remote

Riasztás beállítása hasonló állásokhoz

Principal AI Scientist (m/f/x), Hungary

Ki ☒

 **Egyszerű
jelentkezés**

Mentés

Requirements

As a senior member of a small team, you will have a great influence on the design of SSP's products. Accordingly, we are looking for someone with hands-on experience building production machine-learning solutions from the ground up. In particular, your experience should include typical data science tasks (e.g. idea generation, data manipulation, modeling, model comparison, etc.) as well as basic engineering (e.g. writing clear maintainable code, building data or feature engineering pipelines, maintaining scripts and tooling to support ML tasks and deployments, performing predictions at scale, etc.).

We Appreciate If You Have The Following Experiences

- Strong background in machine learning and related programming in Python (required)
- Hands-on experience productionizing machine learning models/pipelines and a solid understanding of best practices
- Experience with feature engineering and common machine learning algorithms (e.g. Regression, GBM, ANNs, etc.) including knowledge of hyperparameter tuning, regularization, etc.
- Familiarity with general backtesting strategies as well as time-series-specific techniques
- Strong programming background including collaborative coding with version control, unit, and integration testing
- A deep understanding of the utility and limitations of various model evaluation metrics for both online and offline environments
- Ability to stay up to date with ML techniques and libraries and to acquire any missing technical skills concurrently while executing on projects
- Active use of LLM-powered development tools (e.g. Copilot, Cursor, Claude Code, etc.) as part of your regular workflow, including for code generation, refactoring, testing, and documentation — with the critical judgment to evaluate and validate AI-generated output
- Hands-on experience with Databricks and/or PySpark in a production environment
- Experience integrating LLMs or embeddings into production systems (e.g. RAG, semantic search, embeddings-based recommendations)
- Basic data visualization experience to support explanations of problems/solutions

The Soft Skills That We Think Are Important

- Clearly explain problems, present solutions, and debate strategies in English to both technical and non-technical audiences
- Understand a business problem, then propose and weigh multiple solutions
- Collaborate effectively with Data Scientists, Engineers, and Product Team members

What You Bring

- Core ML: feature engineering, model selection, training, hyperparameter tuning, and evaluation, across classical ML (regression, classification, forecasting, ranking) and LLM-based systems
- Python ML stack: pandas/NumPy, scikit-learn, PyTorch or TensorFlow, plus LLM / RAG frameworks
- Rigorous model evaluation: offline metrics, online outcomes, regression suites, faithfulness checks
- MLOps: CI/CD for ML, model and data versioning, monitoring, drift detection, retraining
- Strong SQL and data wrangling on large, messy operational datasets
- Observability tooling: Jenkins, GitHub Actions, Elasticsearch, Kibana
- Cloud-native deployment and REST/OData integration; SAP BTP a plus
- Domain exposure to CPQ, quoting, pricing, or sales/revenue processes a plus

Where you belong

You'll join the team building the customer SAP

Jelentkezés 

Mentés



Qualifications

Apply, If You Have

- 3+ years of experience in Data Science / AI
- Hands-on experience in Python and data preparation techniques (additional R/SQL knowledge is an advantage)
- Understand Large Language Models (LLMs), Retrieval Augmented Generation (RAG) and the transformer architecture (graph RAG is a plus)
- Understand MCPs, comfortably handle model usage monitoring and token utilization concepts in LLM applications
- Demonstrated knowledge of a version control tool (preferably git)
- Deep understanding of data science and machine learning frameworks
- Ability to understand business requirements and translate them into technical requirements
- Experience working with cloud-based technologies
- Experience with Generative AI technology platform offerings (open-source Lang chain/Llama Index, or other 3rd party platforms) and tools
- Hands-on NLP technology experience is a plus
- Team player attitude
- Well-developed written and oral communication skills both in English and in Hungarian

Az álláshirdető által megadott követelmények

- 3+ év munkatapasztalat: Artificial Intelligence (AI)

Riasztás beállítása hasonló állásokhoz

AI Engineer, Budapest, Hungary

Ki ☐



Egyszerű
jelentkezés

Mentés

What We Are Looking For

- 2+ years ML engineering or applied data science experience
- Strong Python; hands-on experience with modern training and fine-tuning workflows
- Experience with structured model evaluation (not just accuracy, precision/recall tradeoffs, drift, edge-case analysis)
- Comfortable working against defined acceptance criteria and production constraints

What You'll Need

- Six or more years in applied machine learning, machine learning engineering, or applied research, with ownership of models shipped into real products
- Depth in at least one of: sequence modeling, latent-variable or probabilistic modeling, temporal modeling, Bayesian methods, or calibration of model outputs, applied to data that changes over time
- Strong Python and production engineering skills: you write the pipelines that train and score your models, and you've shipped models that run on a schedule and serve predictions to real users
- Strong evaluation instincts around calibration, uncertainty, stability, fairness, interpretability, and validation strategy

It Would Be a Bonus If You Had

- Experience with recommender systems, user-state modeling, or personalization at scale
- Experience with knowledge tracing, psychometrics, educational measurement, or adaptive learning systems
- Experience combining structured knowledge representations, such as skills, standards, or concept graphs, with learner models
- Experience designing experiments or observational validation strategies to test whether a model reflects reality