

WHAT DO YOU NEED TO APPLY TO THIS ROLE?

- 4+ years of relevant professional experience, including client-facing or advisory roles
- Advanced Power BI expertise — semantic modeling, DAX, report development, deployment pipelines; you are equally comfortable designing a solution and building it yourself
- Working knowledge of data extraction and processing — SQL, dbt, or Spark; comfortable navigating real-world data environments
- Practical use of AI tools — hands-on experience using LLM-based tools (Claude, Gemini, ChatGPT, or similar) to enhance the speed and quality of your analytical and consulting work
- Strong business acumen — you understand how businesses operate, how decisions get made, and how to frame data in terms of commercial impact
- Exceptional communication skills — able to lead high-level client discussions on data analytics solutions, translate business needs into technical requirements, and present findings with confidence
- Structured, analytical thinking — you approach complex, ambiguous problems with a clear framework and an eye for what matters
- Client ownership mindset — you take responsibility for outcomes, manage expectations proactively, and build long-term trust
- Fluent English, written and spoken
- Degree in Business, Economics, Engineering, Statistics, Computer Science, or a related field

What is good to have?

- SME-level domain knowledge in one or more of: Retail, Banking & Finance, or Telecommunications
- Background in Big4, top-tier consultancy, or CxO-level support in a large enterprise
- Prior experience with generative AI or machine learning projects
- Familiarity with data management methodologies and/or tools
- Proficiency in additional BI tools: Tableau, Looker, or similar platforms

Additional Information

“Data-driven digitalization, human-centered culture”

At Hiflylabs, we strive to create a work environment that is both challenging and supportive, allowing our employees to grow and excel with our company. We believe that our people are our most valuable assets, and we are committed to invest in their personal and professional development through our mentoring system.

Responsibilities

- Design, develop, and deploy innovative Generative AI and Agentic AI solutions from proof of concept to full enterprise production.
- Build and maintain robust pipelines for LLM-based applications, including data processing and retrieval-augmented generation (RAG).
- Implement AgentOps practices to automate the lifecycle of AI agents and LLM applications.
- Collaborate with Product Owners to translate functional requirements into robust technical specifications.
- Ensure the reliability, observability, and compliance of deployed AI solutions through robust monitoring and tracing.
- Guide and mentor junior colleagues to foster AI engineering excellence across teams (SENIOR ONLY)

What You Bring

- Degree in computer science or a similar field.
- 4+ Years in Software Engineering
- Proven hands-on experience in developing and deploying LLM applications.
- Advanced coding skills in Python (preferred) or Java, with a solid understanding of software engineering principles.
- Experience with backend frameworks (e.g., FastAPI, Flask, Django) and building microservice architectures.
- Practical experience with generative AI frameworks like LangChain
- Familiarity with AgentOps tools such as LangSmith, LangFuse, or MLflow for tracing and observability.
- Experience implementing CI/CD pipelines in cloud environments (e.g., Azure, AWS) using tools like Docker, Kubernetes, and GitHub Actions.
- Understanding of Agile, SAFe methodologies and collaborative engineering practices (SENIOR ONLY)
- Fluent English and a proactive, problem-solving mindset.

Nice To Have But Not a Must

- Familiarity with SAP AI Core, SAP BTP
- Experience with agile frameworks (e.g., Scrum, SAFe) and project management tools (e.g., Jira).
- Experience with Agentic AI Frameworks (e.g., LangGraph, CrewAI, Microsoft Agent Framework)

These skills are a plus. If you don't have experience in all these areas, we'll provide the resources and training you need to become a leader in cutting-edge AI technologies.

Where You Beong

CFA is central to making SAP a truly data-driven

Jelentkezés 

Mentés

across platforms

What You'll Bring To The Role

- Strong foundation in computer science or related field (data structures, algorithms, OOP), with hands-on programming experience in languages such as Python or C++ – language is flexible; we value strong coding fundamentals and a willingness to actively code in day-to-day work
- Proven experience in data engineering and distributed computing (5+ years preferred)
- Strong hands-on expertise with Snowflake is useful; experience with other technologies such as Spark, Hadoop, or AWS is advantageous
- Experience delivering complex, large-scale data solutions in a fast-paced environment
- Demonstrated ability to manage multiple priorities and drive end-to-end delivery
- Experience mentoring developers or leading small-to-medium sized teams

Why join us

- Work on strategic, business-critical data platforms

Jelentkezés 

Mentés



Your Ideal Background

- Bachelor's or Master's degree in Engineering, Statistics, Computer Science, Business Administration, or a related field, combined with strong analytical and problem-solving capabilities.
- 5+ years of experience in Business Intelligence, Data Analytics, Reporting, or a related business analytics role.
- Excellent stakeholder management, communication, and project delivery skills, with experience working in Agile and/or Waterfall environments;
- Experience in and Power BI, with a proven track record of designing and delivering executive-level dashboards and reports.
- Some SQL experience and database management experience
- Experience working with Quality, Manufacturing, Supply Chain, Finance, or other operational KPI environments, translating data into actionable business insights.
- Hands-on experience with SAP data and processes, alongside strong proficiency in Microsoft Excel, PowerPoint, and Word.

#jobsEMEAEG

About Te Connectivity

TE Connectivity plc (NYSE: TEL) is a global industrial technology leader creating a safer, sustainable, productive, and connected future. As a trusted innovation partner, our broad range of connectivity and sensor solutions enable the distribution of power, signal and data to advance next-generation transportation, energy networks, automated factories, data centers enabling artificial

Jelentkezés [↗](#)

Mentés

What You'll Bring

- Degree in a technical or quantitative discipline from a top tier institution.
- 3 to 8 years of working experience in software engineering.
- Demonstrated ability to program in Python (preferred) or C++ on Linux/Unix platforms, familiarity with scripting languages.
- Experience with effective code version control, as well as Continuous Integration and Deployment.
- Exceptional communication skills in both verbal and written form.
- Excellent problem-solving abilities and judgment with strong attention to detail.
- Mature and thoughtful, with the ability to operate in a collaborative, team-oriented culture.
- Motivated by the transformational effects of technology-at-scale.
- Experience with AI technologies, such as LLMs and agentic AI, along with their use cases would be beneficial.
- Experience in finance is a plus but not required.

What We Offer

- Competitive compensation package
- Core benefits include: premium private health insurance and life insurance with savings plan
- Support for every aspect of life through Employee Assistance Program and fully covered sick leave
- Strong culture of learning and development: training courses, library, guest speakers, share and learn events, global conferences
- Regular offsite team buildings, annual conferences and occasional global summits – opportunity to travel and connect with our local

[Jelentkezés](#) 

[Mentés](#)

Key Responsibilities

- Develop Generative AI and Agentic AI solutions aligned with Lenovo's strategy and technical standards
- Design and build AI architectures, inference engines, agents, and user interfaces
- Deliver end-to-end AI systems using LLMs, SLMs, VLMs, and relevant frameworks
- Integrate AI solutions into existing platforms, including cloud, on-premise, and edge environments
- Select and implement AI tools and technologies to improve efficiency and scalability
- Ensure security standards (encryption, access control, vulnerability management) across AI solutions
- Apply AIOps and AI development best practices for both experimentation and production
- Create technical documentation and support pre-sales activities (RFPs, solution design)

Requirements

- Bachelor's or Master's degree in Computer Science, AI, or related field
- 5+ years of experience in AI/ML, data science, NLP, computer vision, or Generative AI
- 2+ years of hands-on experience with end-to-end Generative AI solutions (multi-modal preferred)
- Experience with LLM ecosystems (e.g., OpenAI, Anthropic, Llama, Hugging Face, Bedrock)
- Familiarity with NVIDIA technologies (GPUs, CUDA, TensorRT) is an advantage
- Strong understanding of cloud and on-prem environments, including Kubernetes/ containerization
- Proven experience with MLOps / LLMOps / AIOps and CI/CD pipelines (e.g., Jenkins, GitLab)
- Strong programming skills in Python, React, and scripting languages

What We Can Offer You

- An international team with a high focus on Gender Diversity.
- Employee Assistance Program, e.g., for psychological, legal financial consultancy
- You are joining a company that prioritizes sustainable solutions like CO2 Offset, Asset Recovery Services, and the Lenovo Certified Refurbished portfolio.
- Access to training for personal development - Internal E-learning Development Platform Available for Employees
- Mentorship program
- ...and many more!

If you are passionate about AI development and eager to take on a challenging yet rewarding role

You have:

- Degree in a STEM discipline (e.g. Computer Science, Mathematics, Statistics, Engineering, Physics) or equivalent quantitative background
- Experience using SQL, Python, Alteryx or similar to build complex queries for extracting and analyzing data
- Familiarity with database design concepts
- Ability to translate complex technical data concepts into business context
- Working understanding of data protection and data sensitivity rules and best practices
- Excellent written and verbal communication skills
- Experience in Internal Audit, banking or capital markets would be an advantage
- Knowledge of AI/ML concepts and model development would be an advantage
- Experience developing dashboards in Tableau, Power BI or similar would be an advantage
- Practical exposure to GenAI tools and prompt engineering, with an understanding of their limitations, output validation and appropriate use in a controlled environment

Team profile:

The IAD Data Analytics team develops innovative data, automation, and AI solutions that help Internal Audit perform more effective, risk-focused, and scalable assurance activities.

By combining advanced analytics and data engineering capabilities with close partnership across audit and technology teams, the team delivers actionable insights, continuous monitoring capabilities, and forward-looking risk coverage across the Firm.

About us:

Morgan Stanley provides a superior foundation for building a professional career - a place for people to learn, achieve and grow. You will be exposed to a truly international and multi-cultural environment that appreciates and respects individuality. Our state-of-the-art offices in the City Centre have been designed to maximize collaboration. There are various employee offers such as discount theatre tickets, dry cleaning service and many more.



Required Qualifications

- Strong software engineering skills in .NET and/or Python, with experience building production-grade APIs and services
- Applied algorithmic problem solving, including search, ranking, clustering, and data extraction/transformation pipelines
- Solid understanding of LLM architectures and trade-offs (e.g. GPT-family, Claude, Llama, and similar models)
- Experience integrating LLM and agent APIs (e.g. Azure OpenAI, OpenAI API, Claude API, Amazon Bedrock)
- Design and implementation of agent-based systems for automation, orchestration, and decision support
- Experience working with relational, document, and graph databases, including data modeling and querying
- Web data ingestion experience, including web scraping, schema extraction, and handling changing source structures
- Experience building reliable, scalable, maintainable AI systems, including logging, error handling, and performance considerations

Nice To Have

- Fine-tuning and preference optimization techniques (e.g. supervised fine-tuning, RL-style approaches via APIs)
- Designing pipelines that convert unstructured data into reliable structured representations
- Exposure to event-driven or background processing architectures for long-running AI workflows
- Experience contributing to or operating AI-powered systems in production

Riasztás beállítása hasonló állásokhoz

Senior AI Engineer, Budapest, Budapest, Hungary

Ki ☒

Gyorsabb álláskeresés a Premium csomaggal

Hozzáférhet a vállalati információkhoz, például

Jelentkezés

Mentés

Qualifications

- Education: Bachelor's or Master's Degree in Computer Science, Data Science, Software Engineering, or a related quantitative field is required. Solid foundational knowledge in Machine Learning and Software Architecture is essential.
- Experience: 5+ years of significant experience in designing end-to-end technical architectures for Machine Learning and Generative AI solutions. Proven track record in conducting technical feasibility assessments, rapid prototyping (POCs/ MVPs), and bridging business requirements with engineering execution. Experience with enterprise-scale systems, security standards, and AI governance is critical.
- Skills: Expertise in AI design patterns (such as RAG, Fine-tuning, and Agentic workflows) and model evaluation frameworks. Strong technical leadership and mentorship abilities, exceptional stakeholder management across cross-functional teams (Security, Data, Infrastructure), and the capability to optimize architectures for scalability, cost-efficiency, and reliability.
- Languages: Professional proficiency in English required; Italian and/or additional European languages are a plus.

Additional Information

With Over 5,000 Employees Based In More Than 80 Countries Across The Globe, Our Teams Provide Industry Leading Expertise That Spans

- Consumer Engagement
- Parts, Accessories & Service Performance
- Actionable Insights
- Repair Optimization & Compliance

[Jelentkezés](#) 

[Mentés](#)

Your Background

You don't need to tick every box, but we're generally looking for:

- Strong experience in software engineering, technical delivery, or technical leadership
- Experience deploying ML systems into production, including serving, monitoring, and retraining
- Experience leading or mentoring engineering squads
- Hands-on experience building and training deep learning models in PyTorch
- Familiarity with MLOps tooling such as MLflow and CI/CD workflows
- Experience with AWS, Azure, or GCP
- Experience with sequence models (e.g. transformers, RNNs) or generative models (e.g. diffusion, flow, or GANs) – applied to any domain; experience with spatiotemporal, time-series, or motion/trajectory data is a strong plus.
- Experience balancing leadership responsibilities with active development
- Strong stakeholder communication and delivery planning skills
- Willingness to work on-site in our Alkmaar office

Nice to Have (But Not Required)

- Strong understanding of scalable systems, APIs, and cloud-native architectures
- Understanding of real-time data systems or 3D graphics pipelines
- Experience with Unity or game engines
- Interest in football or other team sports

Why Level Up Your Career at Beyond Sports

Jelentkezés 

Mentés



Essential Experience And Skills

- Hands-on experience in AI engineering, building production LLM-powered systems, for example AI-powered chatbots or retrieval-based recommendation tools
- Deep, practical experience building RAG systems from scratch, including chunking strategy, embedding model selection, and indexing design
- Solid understanding of hybrid retrieval: semantic search, reranking, and multi-stage retrieval pipelines
- Experience with vector databases for semantic retrieval at scale
- Experience building evaluation processes that measure retrieval quality against ground truth
- Strong prompt engineering and context engineering skills
- Strong software engineering fundamentals: Python, REST APIs, CI/CD pipelines
- Comfortable working in a fast-moving scale-up environment, figuring things out as you go

Desirable Skills

- Experience with graph databases, as a complement to vector-based retrieval
- Experience with agentic workflow patterns: multi-agent orchestration, tool-calling, state management
- Experience building shared AI infrastructure used by multiple product teams

What We Do

We enable every SAP-run business to manage change intelligently, no matter how complex their technology landscape. Our suite of Intelligent Change Management solutions harnesses the collective intelligence of the SAP community to help business and technology change teams work together to explore, plan, and execute business change imperatives. For over 25 years, we have helped equip, liberate, and champion SAP change

Requirements

- Bachelor's degree or equivalent in Computer Science, Data Science or a related discipline
- At least 2 years of work experience in data and technology fields
- Solid experience in building data-driven applications with a user-focused mindset and visual storytelling with data
- Expertise in Data Analytics and Visualization techniques
- Proficiency in modern programming languages, such as Python or TypeScript
- Good understanding of SQL and relational database concepts
- Experience with Palantir Foundry is a must
- Familiarity with Scrum and other Agile development approaches
- Strong communication skills in English with the ability to translate technical concepts into business-friendly language

Nice to have

- Familiarity with HTML, CSS, JavaScript and Gradle
- Knowledge of the Insurance Domain, Financial Industry or finance-related functions in other industries
- Knowledge of Microsoft Power BI or Tableau
- Skills in PySpark or distributed data processing frameworks
- Experience with Machine Learning or Generative AI models

We offer

- Dynamic, entrepreneurial corporate environment
- Diverse multicultural, multi-functional, and

[Jelentkezés](#) 

[Mentés](#)

Qualifications

- Bachelor's or master's degree in AI, computer science, Engineering, mathematics or a related field.
- 2-4 years of experience in developing and deploying Agentic solutions to production.
- Hands-on experience developing solutions using no code/low code platforms and SDKs.
- Must be proficient in Microsoft Azure GenAI platforms (AI Foundry, Copilot Studio), Power Platform.
- Experience applying advanced AI concept such as MCP servers for connecting AI Agents to tools and.
- Excellent problem-solving skills and attention to detail.
- Strong communication and collaboration skills.

Preferred Skills

- Automation of complex processes using GenAI agents, especially using the Azure Foundry, Copilot Studio, Power Automate and other low code/no code platforms.
- Advanced programming with Python, SQL
- Hands-on experience with at least one of the following data storage systems: Snowflake, Azure Fabric (One Lake), and PostgreSQL
- A preferred candidate would also have experience building sites with SharePoint building automation workflows with Power Automate platform.

Why Join Greif

- Opportunity to work with a global leader in industrial packaging.
- Collaborative and inclusive work environment

Jelentkezés 

Mentés

What will make you successful?

- Bachelor's or Master's degree in Computer Science, AI, Data Science, Engineering, or a related field (or equivalent experience).
- 7+ years of experience in software engineering, AI, cloud, data engineering, or related technical domains, including 3+ years in a technical leadership role.
- Strong programming frontend and backend development skills, preferably in Python, with experience building APIs and scalable applications.
- Hands-on experience with Generative AI technologies, including LLMs, RAG, prompt engineering, vector search, and AI-powered workflows.
- Experience developing and deploying cloud-native solutions, preferably on Microsoft Azure and Azure AI services.
- Solid understanding of software engineering best practices, including testing, CI/CD, Git, observability, security, and production support.
- Knowledge of MLOps/LLMOps, data engineering fundamentals, and secure application integration.
- Proven ability to deliver secure, scalable, production-ready solutions in Agile, cross-functional environments.
- Strong analytical, communication, collaboration, and mentoring skills.
- Passion for innovation and emerging AI technologies.
- Fluent English.

What we offer you

- Working at the world's only fully integrated aluminum and leading renewable energy company
- Diverse, global teams
- Flexible work environment/home office
- We provide you the freedom to be creative and to learn from experts
- Possibility to grow with the company, gain new certificates
- Attractive benefit package

Equal opportunities

Where You Make An Impact

- Design, build, and operate AI agents (single- and multi-agent) including tool use/function calling, orchestration, and state/memory management
- Build and optimize RAG solutions (retrieval, chunking, re-ranking, citations/source grounding, document lifecycle management)
- Implement guardrails and security controls: prompt-injection defenses, PII/DLP handling, policy enforcement, secrets management
- Establish evaluation and testing for agents and GenAI components
- Drive LLMops/observability: logging/tracing, monitoring, feedback loops, A/B testing, cost/token and latency optimization, rollback strategies
- Administer and govern enterprise AI platforms.
- Set up and maintain SSO/access control, role/permission models, provisioning (joiner/mover/leaver), auditing and reporting
- Define and enforce AI governance standards (data classification, approved use cases, retention, compliance/security requirements) together with IT/Security/Legal
- Integrate agents into business workflows via automation/orchestration tools: triggers, workflows, API connectors, error handling, monitoring, retries

Because You Matter (example)

- Competitive salary and benefits package.
- Flexible working hours with a hybrid model require 60% office presence.
- Cafeteria with multiple options in the amount of gross 40,000 HUF
- Private health insurance after 3 months' probation
- Modern office location with ergonomic workspaces, and a collaborative atmosphere

What Makes You Be The Perfect Fit

- Degree in Computer Science, Data Science, Mathematics, Physics, or a related field
- Software engineering skills (e.g., Python), API design/integration, testing, clean code
- Proven experience in agent building (tool/function calling, orchestration patterns, stateful workflows)
- Hands-on expertise in RAG & search, including citations/source-grounding logic

Jelentkezés [↗](#)

Mentés

Where You Make An Impact

- Design, build, and operate AI agents (single- and multi-agent) including tool use/function calling, orchestration, and state/memory management
- Build and optimize RAG solutions (retrieval, chunking, re-ranking, citations/source grounding, document lifecycle management)
- Implement guardrails and security controls: prompt-injection defenses, PII/DLP handling, policy enforcement, secrets management
- Establish evaluation and testing for agents and GenAI components
- Drive LLMOps/observability: logging/tracing, monitoring, feedback loops, A/B testing, cost/token and latency optimization, rollback strategies
- Administer and govern enterprise AI platforms.
- Set up and maintain SSO/access control, role/permission models, provisioning (joiner/mover/leaver), auditing and reporting
- Define and enforce AI governance standards (data classification, approved use cases, retention, compliance/security requirements) together with IT/Security/Legal
- Integrate agents into business workflows via automation/orchestration tools: triggers, workflows, API connectors, error handling, monitoring, retries

Because You Matter (example)

- Competitive salary and benefits package.
- Flexible working hours with a hybrid model require 60% office presence.
- Cafeteria with multiple options in the amount of gross 40,000 HUF
- Private health insurance after 3 months' probation
- Modern office location with ergonomic workspaces, and a collaborative atmosphere

What Makes You Be The Perfect Fit

- Degree in Computer Science, Data Science, Mathematics, Physics, or a related field
- Software engineering skills (e.g., Python), API design/integration, testing, clean code
- Proven experience in agent building (tool/function calling, orchestration patterns, stateful workflows)
- Hands-on expertise in RAG & search, including citations/source-grounding logic
- Experience with administration and governance of multiple enterprise AI platforms, including M365 Copilot, ChatGPT, Gemini, Claude, and Perplexity
- Solid understanding of GenAI security/compliance (PII/DLP, access controls, prompt injection, audit logs)
- Familiarity with system architectures, APIs, and integrating ML models into production environments
- Knowledge of cloud technologies (e.g., Azure, AWS, GCP) and CI/CD processes is a plus
- Analytical mindset, strong team player, and a solution-oriented working style

Ready To Join Us?

Then apply now and become part of the TAKKT family. Together we create impactful, efficient and sustainable solutions designed to improve and simplify our customers' workplace experience – for our customers, for us, and for YOU – **shape new worlds of work!**

Az álláslehetőségek megtekintéséhez kattintson ide.

Jelentkezés [↗](#)

Mentés

Your Skills And Experience That Will Help You Excel

- 4–6 years of experience in roles that combine technology implementation with business/operational problem-solving – solutions engineering, forward deployed engineering, technical consulting, implementation engineering, or similar hybrid roles.
- Bachelor's degree in Computer Science, Data Science, Financial or Quantitative Engineering, or a related field – or equivalent practical experience.
- Strong Python engineering skills, with SQL and PL/SQL for data work and day-to-day Linux/Unix and Git; JavaScript/TypeScript a plus. You write and debug clean, production-quality code.
- Data pipeline, orchestration, and database expertise – workflow engines (Airflow, Prefect), ETL/ELT, SQL and relational databases (Oracle, SQL Server, Azure SQL, Postgres), and big-data/cloud storage (ADLS, Snowflake).
- Experience building market-data systems in a multi-asset-class environment – reference data, corporate actions, benchmark/index, and pricing; familiarity with major data vendors (LSEG/Refinitiv, Bloomberg, S&P) and with financial, ESG/sustainability, or capital-markets data is an advantage.
- Comfortable delivering production-quality software with AI coding assistants (Claude Code, Cursor, GitHub Copilot, or Codex) and working in Agile/DevOps teams with CI/CD and automated testing; familiarity with MCP or reusable Skills is a plus.
- Hands-on experience building AI/ML-powered solutions – agent frameworks (LangGraph, CrewAI, or AutoGen), LLM APIs (Anthropic Claude, OpenAI), prompt engineering, RAG, and evaluation – and deploying and monitoring agentic systems at scale, are strong advantages.
- Demonstrated success shipping products quickly inside a tightly governed enterprise – meeting rigorous security, compliance, and change-control standards without sacrificing delivery speed.
- Systems thinking and strong communication – you design coherent end-to-end systems and explain technical decisions clearly to both engineers and non-technical stakeholders.

About MSCI

What we offer you

- Salary range: HUF 16,150,000 - HUF 17,765,000 / year plus eligible for annual bonus
- Transparent compensation schemes and comprehensive employee benefits, tailored to your location, ensuring your financial security, health, and overall wellbeing.
- Flexible working arrangements, advanced technology, and collaborative workspaces.
- A culture of high performance and innovation

Jelentkezés 

Mentés

Huge plus if you

- Have any experience with or willingness to learn the following technologies: distributed systems, NodeJS, Javascript/Typescript, Go, and AWS.
- Worked with any of these technologies: Amplitude, Apify, Builtwith, [Chart.io](#) Dropcontact, Dux-Soup, HubSpot, [Hunter.io](#), Intercom, Pardot, Drift, Marketo, Integromat, Lemlist, N8N, [Outreach.io](#), [Reply.io](#), Segment, Woodpecker, Zapier, or similar.

About you

- You are deeply excited about growing a startup using data and tech
- Interested in building streamlined automation, scripts, and scrapers and architecting solutions that help business teams simplify, automate, and grow - with the ability to work across the entire stack
- Understand data flows and databases and the tools that power marketing automation, experimentation, personalization, attribution, and analytics
- You're good with SQL (we use Postgres) and you want to learn some MongoDB, and Python as well
- You understand what is API economy and you like to work with APIs and get the most out of the data they provide
- Proactive self-starter with the ability to be hands-on and can analyze, interpret and derive insights into business reports, analytic models, and data sources
- Sharp technical troubleshooting and ability to collaborate with business teams
- A quick learner that values making data-driven recommendations
- You have a basic understanding and experience with cloud providers (we use AWS)

Why Recart?

- A product loved by thousands of Shopify stores across the globe
- Above-average stock options (we really mean ownership)
- Flexible vacation days and working hours
- Clear quarterly OKRs and feedback loops. No chaos, no micro-management
- A team that works hard and supports each other with empathy

Essential technical experience:

- Strong Python/software engineering background
- Commercial experience building **production**

LLM-powered products or systems

- Deep hands-on RAG/retrieval engineering experience
- Vector databases and semantic search
- Hybrid and multi-stage retrieval
- Reranking techniques
- Retrieval evaluation against ground truth
- Prompt and context engineering
- REST APIs
- CI/CD and production deployment

Experience with any of the following would be particularly valuable:

- Graph databases / GraphRAG
- LangGraph or similar orchestration frameworks
- Multi-agent systems
- Tool calling
- MCP
- State management
- AI observability/evaluation platforms
- Shared AI infrastructure or AI platforms

We're particularly interested in engineers who have worked within **product-led technology environments**, where you've owned AI capabilities beyond the initial prototype and seen how they behave with real users and real production data. You might currently be working for an **AI product company, SaaS business or product engineering organisation**, or you may have come from consulting and now want the opportunity to own one AI product deeply rather than moving between short-term customer projects.

Either way, we're looking for someone who thinks like a **product engineer**: build it, measure it, understand why it fails, improve it and own the

Jelentkezés [↗](#)

Mentés

About You

- 4-7 years of experience in marketing analytics, marketing effectiveness, media analytics, consumer insights or other analytical roles.
- Strong understanding of marketing measurement methodologies, including MMM, incrementality testing, ROI measurement and media effectiveness.
- Proven ability to translate complex data and analysis into clear business recommendations.
- Experience working with large datasets, multiple data sources and performance metrics.
- Strong stakeholder management and communication skills, with the ability to influence non-technical audiences.
- Experience managing multiple projects and priorities in a fast-paced environment.
- Strong analytical, problem-solving and data storytelling capabilities.
- Advanced Excel and PowerPoint skills, with familiarity with marketing effectiveness and analytics tools.
- Curious, proactive and eager to learn new methodologies, technologies and ways of working.

Benefits

- Unique Family Leave opportunity
- We work without border - you will work with people from all over the world
- You will deal with world class brand portfolio
- Wide variety of thematic events, afterwork activities organized by employee communities
- Product Allowance to purchase company products from the staff shop
- We won the Fittest Workplace Award, of course we have a Gym
- Urban legend is true, we do have a Rooftop Bar

[Jelentkezés](#) 

[Mentés](#)

Required Qualifications:

- 8+ years working as a Data Engineer, Analytics Engineer, or Data Analyst in a data-driven environment (Staff level typically 10+ or equivalent scope/impact).
- Strong SQL skills (complex joins, window functions, aggregation, query optimization).
- Experience with at least one major data warehouse / big data technology (e.g., BigQuery, Snowflake, Redshift, Hive, Spark, Vertica, StarRocks, or similar).
- Solid experience with ETL/ELT tools or orchestration frameworks (e.g., Airflow, dbt, internal frameworks).
- Proficiency in at least one programming language for data and services (e.g., Python, Java, Golang, Kotlin, or JavaScript/TypeScript).
- Experience building production-grade data pipelines and/or data services with attention to quality, monitoring, and maintainability.
- Experience with BI / visualization tools (e.g., Tableau, Power BI, Superset, Metabase, Looker).
- Strong analytical thinking: comfortable framing business questions, exploring data, and turning results into clear insights.
- Ability to communicate complex topics in simple, concise language to non-technical stakeholders.

Preferred Qualifications:

- Experience with event-driven data (logs, clickstream, customer journey, contact flows, tracking data).
- Knowledge of data modeling techniques (e.g., dimensional modeling, Kimball, data vault, star/snowflake schema).
- Experience with A/B testing, experimentation platforms, or causal analysis.
- Familiarity with ML pipelines or working with data science teams (e.g., routing models, QA scoring, forecasting, recommendations).
- Experience in online consumer products, travel, e-commerce, or marketplaces.
- Hands-on frontend or backend development experience (e.g., building internal tools in React/Next.js, Node.js, or similar frameworks) used by operations or support teams.
- Prior experience in a Staff or Tech-Lead role, driving technical direction and mentoring multiple engineers/analysts.

#sanfrancisco #sanjose #losangeles #sandiego
#oakland #denver #miami #orlando #atlanta
#chicago #boston #detroit #newyork #portland
#philadelphia #dallas #houston #austin #seattle
#sydney #melbourne #perth #toronto #vancouver
#montreal #shanghai #beijing #shenzhen #prague
#Brno #Ostrava #cairo #alexandria #giza #estonia
#paris #berlin #munich #hamburg #stuttgart
#cologne #frankfurt #hongkong #budapest
#jakarta #bali #dublin #telaviv #milan #rome
#venice #florence #naples #turin #palermo

Jelentkezés [↗](#)

Mentés

Requirements

- 6+ years of experience in ML Engineering, MLOps or DevOps for AI/ML systems (or 3+ years with a PhD)
- Expertise in MLOps principles and tools (testing, deployment, monitoring, versioning)
- Proven experience building production-grade ML systems in Python
- Hands-on experience with cloud platforms (AWS preferred) and container orchestration (Docker, Kubernetes)
- Experience with ML platforms such as MLflow, SageMaker, Vertex AI or similar
- Strong understanding of the complete ML lifecycle, including training, validation, retraining, A/B testing and drift detection
- Experience with CI/CD pipelines (GitLab CI, GitHub Actions, Jenkins) and Infrastructure as Code (Terraform, CloudFormation)
- Proficiency in Python and experience with scalable data processing frameworks (Spark, Airflow, Kubeflow)
- Familiarity with monitoring/logging tools like Prometheus, Grafana and Datadog
- Professional English proficiency
- Ability to work from Paris, Milan, Turin, Barcelona or remotely within a compatible time zone

Nice to have

- Experience deploying or scaling GenAI solutions
- Contributions to shared internal libraries or developer tools
- Additional language skills (French, Spanish, Italian)
- Familiarity with large-scale distributed systems or marketplace environments

Jelentkezés 

Mentés

These are the essentials you'll need to get an interview:

- A degree in Computer Science, Software Engineering, or a related IT discipline
- Demonstrable understanding of AI/ML concepts, including LLMs and agent-based systems
- Ability to write clean, well-tested code — Python preferred
- Strong experience of shipping real AI products - not chatbots!
- Strong communication skills and a collaborative mindset

It would be great if you had these to, but we'll support you if you don't:

- Exposure to AWS services, especially Amazon Bedrock or similar LLM hosting platforms
- Familiarity with GRC concepts (governance, risk, compliance) or cybersecurity frameworks such as ISO 27001 or SOC 2
- Any experience with agent orchestration libraries (e.g. LangChain, CrewAI, AutoGen)

About Us

Diligent is the AI leader in governance, risk and compliance (GRC) SaaS solutions, helping more than 1 million users and 700,000 board members to clarify risk and elevate governance. The Diligent One Platform gives practitioners, the C-Suite and the board a consolidated view of their entire GRC practice so they can more effectively manage risk,

[Jelentkezés](#) 

[Mentés](#)

These are the essentials you'll need to get an interview

- Direct experience building mixed-method behavioural insight systems — you've built the layer where quantitative behaviour and qualitative evidence come together, not just consumed the output. You've worked alongside design or UX researchers, and can name a time the qualitative read changed what you measured.
- Hands-on event and telemetry schema design — Pendo, Amplitude, Segment, Mixpanel or equivalent. You designed the schema and wrote the spec, not used dashboards someone else built.
- Designing and governing product data models and taxonomies — canonical metrics and naming taken into a multi-product estate and made to stick, including acquired products.
- Data governance practice: schema and field ownership, versioned change control, standards adoption across teams you don't own.
- Fluency in SQL and the BI / warehouse layer — you build and validate roll-up logic yourself.
- Behavioural segmentation judgement — you know why a B2B account is a mix of people with different motivations, and why averaging them produces confident, wrong conclusions.
- The credibility to influence without authority — you can hold a standard with teams you don't line-manage.

It would be great if you had these too, but we'll support you if you don't

- Hands-on behavioural analytics and synthesis practice — running the analysis yourself, not only building what enables it.
- Privacy and lawful basis as live design constraints — pseudonymisation, residency, re-identification controls.
- GRC, governance or regulated-software domains.
- Salesforce, Jira Product Discovery, Pendo, and consolidating fragmented tooling into one system of record.
- Product value / ROI measurement — connecting behaviour to commercial outcomes.

How We'll Assess This

We ask everyone to write SQL and design an event schema, and to talk us through a time qualitative evidence contradicted their numbers. No take-home longer than an hour.

Here's where this role goes next

We're building this function, not maintaining it. You

Requirements

- 7+ years of engineering experience, with a strong recent track record building production AI / LLM applications (not prototypes or research only)
- Strong agent-design judgment — task-harness fit, matching the harness to the context, failures, and policies of the actual task rather than calling a model in a loop
- The ability to operate close to the client: lead discovery and feasibility conversations, work directly with SMEs and end-users, and explain technical trade-offs to both technical and non-technical audiences
- Hands-on experience with agentic frameworks (LangChain, LangGraph, Semantic Kernel, or similar) and major LLM providers (OpenAI, Anthropic, Google Gemini)
- Expert-level Python and solid software engineering fundamentals
- Strong RAG and retrieval skills: vector databases, embeddings, hybrid search, re-ranking, chunking, and context management
- Proven experience evaluating generative AI quality — LLM-based evaluation, heuristics, custom eval frameworks — and using observability/tracing tools (LangSmith, Arize Phoenix, Langfuse, or similar)
- Production deployment experience on at least one major cloud (AWS, Azure, or GCP) with containerization, CI/CD
- Sound judgment under ambiguity — scoping, sequencing, and making the call on speed vs. quality vs. scope
- English at C1 level

Nice to have

- Experience designing experiments, A/B testing, and iterating on AI products against real user behavior and business metrics
- Background in NLP, Data Science, or applied ML, with experience moving models into production
- Familiarity with MCP, A2A, Agent Skills, and emerging agent standards
- Experience with enterprise AI platforms (AWS Bedrock AgentCore, Databricks Genie, Microsoft Foundry, Gemini Enterprise)
- Exposure to AI governance, security, and compliance (guardrails, prompt-injection prevention)
- Prior client-facing or pre-sales exposure in a consulting or services context

We offer

- Dynamic, entrepreneurial corporate environment
- Diverse multicultural, multi-functional, and multilingual work environment
- Opportunities for personal and career growth in a progressive industry
- Global scope, international projects

[Jelentkezés](#) 

[Mentés](#)

Skills & Qualifications

- Minimum of 3 years' proven experience within financial services or asset management, preferably across multiple asset classes and investment subject areas
- Strong analytical and quantitative skills, with the desire to solve investment problems and translate analysis into practical insights. Demonstrated experience in improving processes, analytics, and workflows through technology. Coding skills and/or knowledge of Aladdin, Morningstar, Bloomberg, Tableau, VBA, SQL, and/or Python is expected, with a willingness to further develop these skills to enhance portfolio construction capabilities.
- Solid understanding of financial markets, investments (i.e. Fixed Income, Equities, Alternatives) and key financial concepts (such as asset allocation, risk analytics, performance, attribution etc.)
- CFA and/or CAIA encouraged but not required.
- Consistent track record of academic achievement, work ethic, and leadership.
- Fluent in English with excellent communication skills, able to communicate effectively with a range of internal and external partners, and able to explain complex topics to non-investment individuals, both verbally and in writing.
- Curious self-starter passionate about execution with results-oriented work ethic and professional disposition.
- Strong teammate who takes a collaborative and cooperative approach to tasks, experienced working with multiple partners to deliver on projects on time and to a high standard, but can work independently.
- Organized and thorough, able to prioritise and multi-task quickly efficiently, adapting to new situations in a fast-paced environment.
- Being a strong fiduciary by ensuring the highest level of integrity and professionalism in every aspect of performance.

The salary range for Hungary is 13400000 - 20100000 HUF

Additionally, employees are eligible for an annual discretionary bonus, and benefits including healthcare, leave benefits, and retirement benefits. BlackRock operates a pay-for-performance compensation philosophy and your total compensation may vary based on role, location, and firm, department and individual performance.

Our Benefits

To help you stay energized, engaged and inspired, we offer a wide range of employee benefits including: retirement investment and tools designed to help you in building a sound financial future; access to education reimbursement; comprehensive resources to support your physical health and emotional well-being; family support programs; and Flexible Time Off (FTO) so you can relax, recharge and be there for the people you care about.

Our hybrid work model

BlackRock's hybrid work model is designed to enable a culture of collaboration and apprenticeship that enriches the experience of our employees, while supporting flexibility for all. Employees are currently required to work at least 4

[Jelentkezés](#)

[Mentés](#)